

## Beyla vs Hubble/Cilium

- **Beyla:** Requires additional services like Tempo, OTEL Collector, Prometheus. Becomes memory-intensive. Good for L7 metrics. L3 metrics can be enabled in Beyla configmap and queried in grafana/prometheus(beyla\_network\_flow\_bytes\_total). Can integrate with grafana to display service graphs in grafana dashboards. Does not need to be deployed before any other services, can dynamically add services to the service discovery and remove at any time based on requirements.
- **Hubble/Cilium** - Service graphs are visible in Hubble UI. To view in grafana, we need to purchase Isovalent Enterprise. Good for L7 metrics and dashboards. L3 metrics can be enabled also, but are not easily exportable to grafana or are not as accessible/detailed in Hubble UI. Less memory-intensive, less external configurations. But needs to be started before any other services, monitors all running services in the cluster.

## Enabling Service Graphs in Grafana using Beyla, Tempo, and OpenTelemetry

### Create a GKE cluster

```
Unset
gcloud container clusters create beyla-test
```

### Create beyla namespace

```
Unset
kubectl create namespace beyla
```

### Create beyla service account and role binding

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: beyla
  namespace: beyla
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: beyla
rules:
- apiGroups: [ "apps" ]
  resources: [ "replicasets" ]
```

```

    verbs: [ "list", "watch" ]
  - apiGroups: [ "" ]
    resources: [ "pods", "services", "nodes" ]
    verbs: [ "list", "watch" ]
  - apiGroups: [ "" ]
    resources: [ "pods", "nodes" ]
    verbs: [ "get", "list", "watch" ]
  - apiGroups: [ "apps" ]
    resources: [ "deployments" ]
    verbs: [ "get", "list", "watch" ]
  - apiGroups: [ "monitoring.coreos.com" ]
    resources: [ "servicemonitors" ]
    verbs: [ "get", "list", "watch", "create" ]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: beyla
subjects:
  - kind: ServiceAccount
    name: beyla
    namespace: beyla
roleRef:
  kind: ClusterRole
  name: beyla
  apiGroup: rbac.authorization.k8s.io

```

## Deploy Beyla as a DaemonSet

beyla-config and beyla daemonset: (Beyla 1.6.4)

**beyla\_network\_flow\_bytes** metric can be enabled using the highlighted section

```

apiVersion: v1
kind: ConfigMap
metadata:
  namespace: beyla
  name: beyla-config
data:
  beyla-config.yml: |
    # this is for beyla_network_flow_bytes metric
    network:
      enable: true
      print_flows: true
      attributes:

```

```

    kubernetes:
      enable: true
      select:
        beyla_network_flow_bytes:
          include:
            - beyla.ip
            - src.name
            - dst.port
            - k8s.src.owner.name
            - k8s.src.namespace
            - k8s.dst.owner.name
            - k8s.dst.namespace
            - k8s.cluster.name

# this will provide automatic routes report while minimizing cardinality
routes:
  unmatched: heuristic
#service discovery
discovery:
  services:
    - k8s_deployment_name: "^docs$"
    - open_ports: 8080-8089
    - k8s_deployment_name: "^website$"
    - k8s_namespace: default
    - k8s_namespace: beyla
    - k8s_deployment_name: "^client-service$"
    - k8s_deployment_name: "^hello-service$"
otel_traces_export:
  reporters_cache_len: 1024
  sampler:
    name: "always_on"
  endpoint: http://opentelemetrycollector.beyla:4317
---
apiVersion: apps/v1
kind: DaemonSet
metadata:
  namespace: beyla
  name: beyla
spec:
  selector:
    matchLabels:
      instrumentation: beyla
  template:

```

```
metadata:
  labels:
    instrumentation: beyla
spec:
  serviceAccountName: beyla
  hostPID: true # mandatory!
  hostNetwork: true
  containers:
    - name: beyla
      image: grafana/beyla:1.6.4
      ports:
        - containerPort: 9092
          hostPort: 9092
          name: http-metrics
          protocol: TCP
      imagePullPolicy: IfNotPresent
      securityContext:
        capabilities:
          add:
            - BPF
            - PERFMON
            - NET_ADMIN
            - SYS_RESOURCE
        privileged: true # mandatory!
        readOnlyRootFilesystem: true
      volumeMounts:
        - mountPath: /config
          name: beyla-config
        - mountPath: /var/run/beyla
          name: var-run-beyla
      env:
        - name: BEYLA_CONFIG_PATH
          value: "/config/beyla-config.yml"
        - name: BEYLA_PROMETHEUS_PORT
          value: "9092"
        - name: BEYLA_PRINT_TRACES
          value: 'true'
        - name: OTEL_EXPORTER_OTLP_ENDPOINT
          value: http://opentelemetrycollector.beyla.svc.cluster.local:4317
  volumes:
    - name: beyla-config
      configMap:
```

```
    name: beyla-config
- name: var-run-beyla
  emptyDir: {}
```

## Deploy OpenTelemetry Collector

### Service:

```
apiVersion: v1
kind: Service
metadata:
  name: opentelemetrycollector
  namespace: beyla
spec:
  ports:
    - name: grpc-otlp
      port: 4317
      protocol: TCP
      targetPort: 4317
  selector:
    app.kubernetes.io/name: opentelemetrycollector
  type: ClusterIP
```

### Configmap:

```
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: collector-config
  namespace: beyla
data:
  collector.yaml: |
    receivers:
      otlp:
        protocols:
          grpc:
            endpoint: ${env:MY_POD_IP}:4317
          http:
            endpoint: ${env:MY_POD_IP}:4318
    processors:
      batch:
      memory_limiter:
```

```
# 80% of maximum memory up to 2G
limit_mib: 1500
# 25% of limit up to 2G
spike_limit_mib: 512
check_interval: 5s
extensions:
  zpages: {}
exporters:
  # logging:
  otlp:
    endpoint: "tempo-distributor.tempo-test:4317"
    insecure: true
service:
  extensions: [zpages]
  pipelines:
    traces/1:
      receivers: [otlp]
      processors: [memory_limiter, batch]
      exporters: [otlp]
    # metrics:
    #   receivers: [otlp]
    #   processors: [batch]
    #   exporters: [logging]
```

## Deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: opentelemetrycollector
  namespace: beyla
spec:
  replicas: 1
  selector:
    matchLabels:
      app.kubernetes.io/name: opentelemetrycollector
  template:
    metadata:
      labels:
        app.kubernetes.io/name: opentelemetrycollector
    spec:
      containers:
```

```

- name: otelcol
  args:
    - --config=/conf/collector.yaml
  image: otel/opentelemetry-collector:0.18.0
  volumeMounts:
    - mountPath: /conf
      name: collector-config
  volumes:
    - configMap:
        items:
          - key: collector.yaml
            path: collector.yaml
            name: collector-config
        name: collector-config

```

## Deploy Client-service and Hello-service to generate sample traces and metrics

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-service
  namespace: beyla
spec:
  replicas: 1
  selector:
    matchLabels:
      app: hello-service
  template:
    metadata:
      labels:
        app: hello-service
    spec:
      containers:
        - name: hello-service
          image: hashicorp/http-echo:0.2.3
          args:
            - "-text=Hello from Hello Service"
          ports:
            - containerPort: 5678
---
apiVersion: apps/v1
kind: Deployment

```

```
metadata:
  name: client-service
  namespace: beyla
spec:
  replicas: 1
  selector:
    matchLabels:
      app: client-service
  template:
    metadata:
      labels:
        app: client-service
    spec:
      containers:
        - name: client-service
          image: curlimages/curl:7.83.1
          args:
            - "/bin/sh"
            - "-c"
            - "while true; do curl -s hello-service.beyla; sleep 5; done"
---
apiVersion: v1
kind: Service
metadata:
  name: hello-service
  namespace: beyla
spec:
  selector:
    app: hello-service
  ports:
    - protocol: TCP
      port: 80
      targetPort: 5678
---
apiVersion: v1
kind: Service
metadata:
  name: client-service
  namespace: beyla
spec:
  selector:
    app: client-service
```



```
ports:
  - protocol: TCP
    port: 80
    targetPort: 80
```

## Deploy Prometheus

1. Create namespace 'prometheus'
2. Prometheus deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: prometheus-server
  namespace: prometheus
spec:
  replicas: 1
  selector:
    matchLabels:
      app: prometheus-server
  template:
    metadata:
      labels:
        app: prometheus-server
    spec:
      containers:
        - name: prometheus
          image: prom/prometheus
          args:
            - "--config.file=/etc/prometheus/prometheus.yml"
            - "--web.enable-remote-write-receiver"
          ports:
            - containerPort: 9090
          volumeMounts:
            - name: config-volume
              mountPath: /etc/prometheus
      volumes:
        - name: config-volume
          configMap:
            name: prometheus-server-conf
            defaultMode: 420
```

3. Prometheus configmap:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: prometheus-server-conf
  namespace: prometheus
data:
  prometheus.yml: |
    global:
      scrape_interval: 15s
      evaluation_interval: 15s
    scrape_configs:
      - job_name: 'beyla'
        static_configs:
          - targets: ['beyla.beyla:9092']
```

#### 4. Prometheus Service:

```
apiVersion: v1
kind: Service
metadata:
  name: prometheus-service
  namespace: prometheus
spec:
  selector:
    app: prometheus-server
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9090
  type: LoadBalancer
```

## Deploy Grafana

1. Create namespace grafana
2. Grafana Deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: grafana
  namespace: grafana
spec:
  replicas: 1
  selector:
```

```

    matchLabels:
      app: grafana
  template:
    metadata:
      labels:
        app: grafana
    spec:
      containers:
        - name: grafana
          image: grafana/grafana
          ports:
            - containerPort: 3000
          env:
            - name: GF_SECURITY_ADMIN_PASSWORD
              value: "admin"
---
apiVersion: v1
kind: Service
metadata:
  name: grafana
  namespace: grafana
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 3000
      protocol: TCP
  selector:
    app: grafana

```

## Deploy Grafana Tempo

1. Create namespace 'tempo-test'
2. Install tempo distributed

Unset

```
helm -n tempo-test install tempo grafana/tempo-distributed
```

3. Edit the helm release Values to enable metrics generator:

Go to the Helm releases in Lens in the tempo-test namespace. Select the tempo release. Edit the Values file:

Add the defaults section to global\_overrides: (highlighted)

```
global_overrides:
```

```

defaults:
  metrics_generator:
    processor:
      service_graphs: null
      span_metrics: null
    processors:
      - service-graphs
      - span-metrics
  per_tenant_override_config: /runtime-config/overrides.yaml

```

Under config: storage: set remote\_write: -url: (highlighted in red)

“<http://prometheus-service.prometheus/api/v1/write>”

Also change metricsGenerator.enabled from false to true. (highlighted)

```

metricsGenerator:
  affinity: |
    podAntiAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        - labelSelector:
            matchLabels:
              {{- include "tempo.selectorLabels" (dict "ctx" . "component"
"metrics-generator") | nindent 10 }}
            topologyKey: kubernetes.io/hostname
        preferredDuringSchedulingIgnoredDuringExecution:
          - weight: 100
            podAffinityTerm:
              labelSelector:
                matchLabels:
                  {{- include "tempo.selectorLabels" (dict "ctx" . "component"
"metrics-generator") | nindent 12 }}
              topologyKey: topology.kubernetes.io/zone
  annotations: {}
  appProtocol:
    grpc: null
  config:
    metrics_ingestion_time_range_slack: 30s
    processor:
      service_graphs:
        dimensions: []
        histogram_buckets:
          - 0.1
          - 0.2
          - 0.4

```

```
- 0.8
- 1.6
- 3.2
- 6.4
- 12.8
max_items: 10000
wait: 10s
workers: 10
span_metrics:
  dimensions: []
  histogram_buckets:
    - 0.002
    - 0.004
    - 0.008
    - 0.016
    - 0.032
    - 0.064
    - 0.128
    - 0.256
    - 0.512
    - 1.02
    - 2.05
    - 4.1
registry:
  collection_interval: 15s
  external_labels: {}
  stale_duration: 15m
storage:
  path: /var/tempo/wal
  remote_write:
    - url: http://prometheus-service.prometheus/api/v1/write
    remote_write_add_org_id_header: true
    remote_write_flush_deadline: 1m
  wal: null
traces_storage:
  path: /var/tempo/traces
enabled: true
extraArgs: []
extraEnv: []
extraEnvFrom: []
extraVolumeMounts: []
extraVolumes: []
```

```
hostAliases: []
image:
  pullSecrets: []
  registry: null
  repository: null
  tag: null
initContainers: []
kind: Deployment
nodeSelector: {}
persistence:
  annotations: {}
  enabled: false
  size: 10Gi
  storageClass: null
podAnnotations: {}
podLabels: {}
ports:
- name: grpc
  port: 9095
  service: true
- name: http-memberlist
  port: 7946
  service: false
- name: http-metrics
  port: 3100
  service: true
priorityClassName: null
replicas: 1
resources: {}
service:
  annotations: {}
terminationGracePeriodSeconds: 300
tolerations: []
topologySpreadConstraints: |
  - maxSkew: 1
    topologyKey: topology.kubernetes.io/zone
    whenUnsatisfiable: ScheduleAnyway
    labelSelector:
      matchLabels:
        {{- include "tempo.selectorLabels" (dict "ctx" . "component"
"metrics-generator") | nindent 6 }}
walEmptyDir: {}
```

Save and deploy a new release of tempo with the edited Values. Metrics Generator pods should now appear in the tempo-test namespace.

## Create prometheus and tempo data sources in Grafana.

Prometheus: <http://prometheus-service.prometheus:80>

Tempo: <http://tempo-query-frontend.tempo-test:3100>

Tempo data source additional settings:

- Trace to metrics data source - prometheus
- Additional settings - service graph data source - prometheus

## View Service Graphs

Go to the Explore tab in Grafana. Click on the service graph tab in the query.

### IF OTEL COLLECTOR OR OTHER PODS KEEP GETTING EVICTED/KEEP RESTARTING DUE TO MEMORY PRESSURE:

- INCREASE NODE POOL SIZE OF THE GKE CLUSTER.
- REMOVE THE BEYLA NAMESPACE WATCH FROM THE BEYLA CONFIGMAP

Change these settings to the following in the tempo helm chart if there is some ResourceExhausted error in the OTEL collector logs when sending to tempo:

Unset

**server:**

**grpc\_server\_max\_recv\_msg\_size:** 123412341234123

**grpc\_server\_max\_send\_msg\_size:** 123412341234123